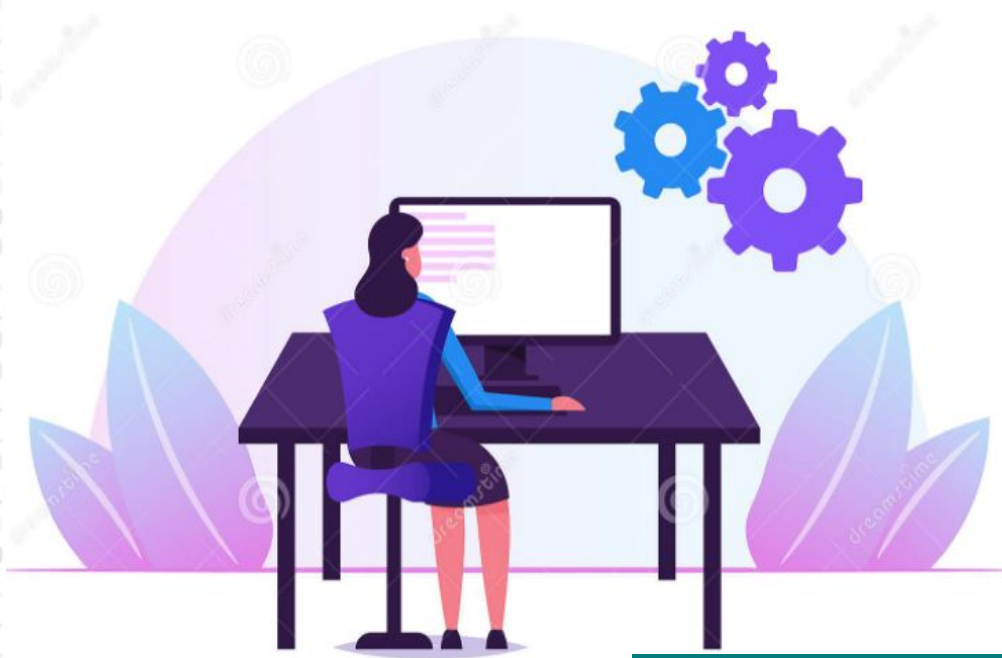


# Welcome!

- We'll start in a moment :)
- We are recording tonight's event. We may plan to take screenshots for social media.
  - ***If you want to remain anonymous***, change your name & keep video off.
- We'll introduce the hosts and break in-between for Q/A.
- We will make some time for Q&A at the end of the presentation as well.
- You can come prepared with questions. And, feel free to take notes.
- Online event best practices:
  - Don't multitask. Distractions reduce your ability to remember concepts.
  - Mute yourself when you aren't talking.
  - We want the session to be interactive.
  - Feel free to unmute and ask questions in the middle of the presentation.
  - Turn on your video if you feel comfortable.
  - Disclaimer: Speaker doesn't know everything!

## Check out:

- [Technical Tracks](#) and [Digital Events](#)
- Get updates – join the [Digital mailing list](#)
- Give us your feedback – take the [Survey](#)



# WWCode Digital + **Backend** **Backend Study Group**

July 15, 2021

Copyright © 2021 by [Prachi Shah](#)

WOMEN WHO  
**CODE**

# Introduction & Agenda

- Welcome from WWCode!
- Our mission: Inspiring women to excel in technology careers.
- Our vision: A world where women are representative as technical executives, founders, VCs, board members and software engineers.
- What is Backend Engineering?
- **Insights into data engineering, data science and machine learning engineering**
  - Data engineering [Part 1 of 2]
    - + Introduction
    - + Similarities/Differences
    - + Day in a life of data engineer
    - + Tech stack
  - Data science and machine learning engineering [Part 2 of 2]



Prachi Shah  
**Senior Software  
Engineer @ Metromile**



Madhurima Nath  
**Data Scientist @  
Slalom**



# Backend Engineering

- What is Backend Engineering?
- Design, build and maintain server-side web applications.
- Concepts: Client-server architecture, API, micro-service, database engineering, distributed systems, storage, performance, deployment, availability, monitoring, etc.

## Software Design

- Defining the architecture, modules, interfaces and data.
- Solve a problem or build a product.
- Define the input, output, business rules, data schema.
- Design patterns solve common problems.
- 3 Types:
  - UI design: Data visualization and presentation.
  - Data design: Data representation and storage.
  - Process design: Validation, manipulation and storage of data.

# What is data engineering?

Data engineers

- **design and build pipelines** to transform and transport data into a format readily useable by the data scientists or other end users.
- pipelines take data from many disparate sources and collect them into a single warehouse that is the unified data source for others.
- work closely with **DevOps and data science/machine learning** teams

# Data Engineer (DE) vs Data Scientist (DS) vs Machine Learning Engineer (MLE)

## **Data engineer:**

builds and develops pipelines, and maintains of data infrastructure, either on-premises or in the cloud (or hybrid or multi-cloud), comprising of databases or data warehouses

## **Data scientist:**

builds and develops mathematical and statistical models -- called machine learning models, to find patterns and gain more insights from the data

## **Machine learning engineer:**

design architecture and pipelines (or software) to integrate and automate the process of running the models developed by data scientists with the entire infrastructure

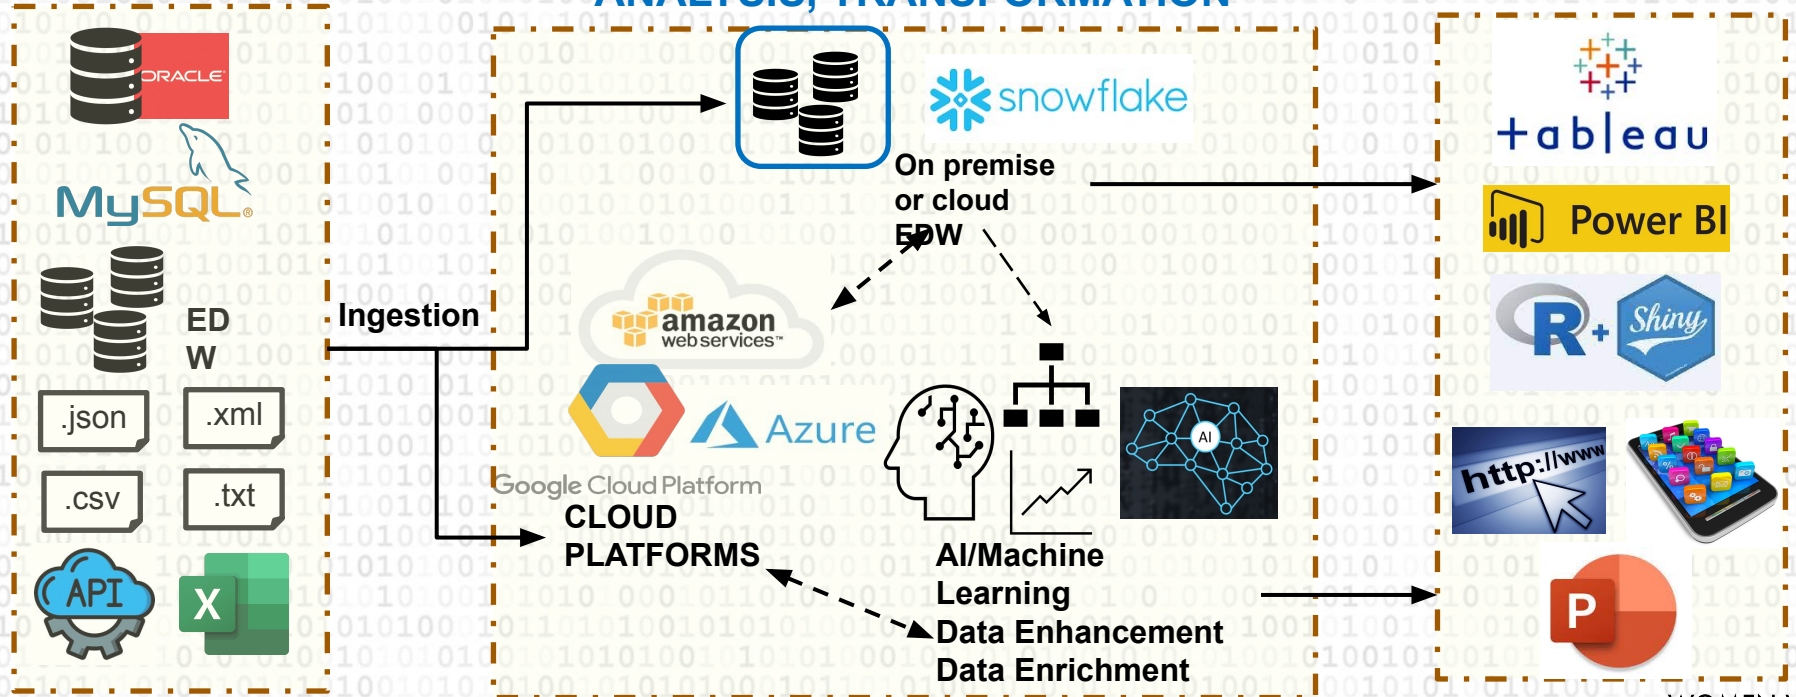


# Data Architecture Diagram

## DATA SOURCES

## DATA CLEANING, PROCESSING, ANALYSIS, TRANSFORMATION

## USER INTERFACE

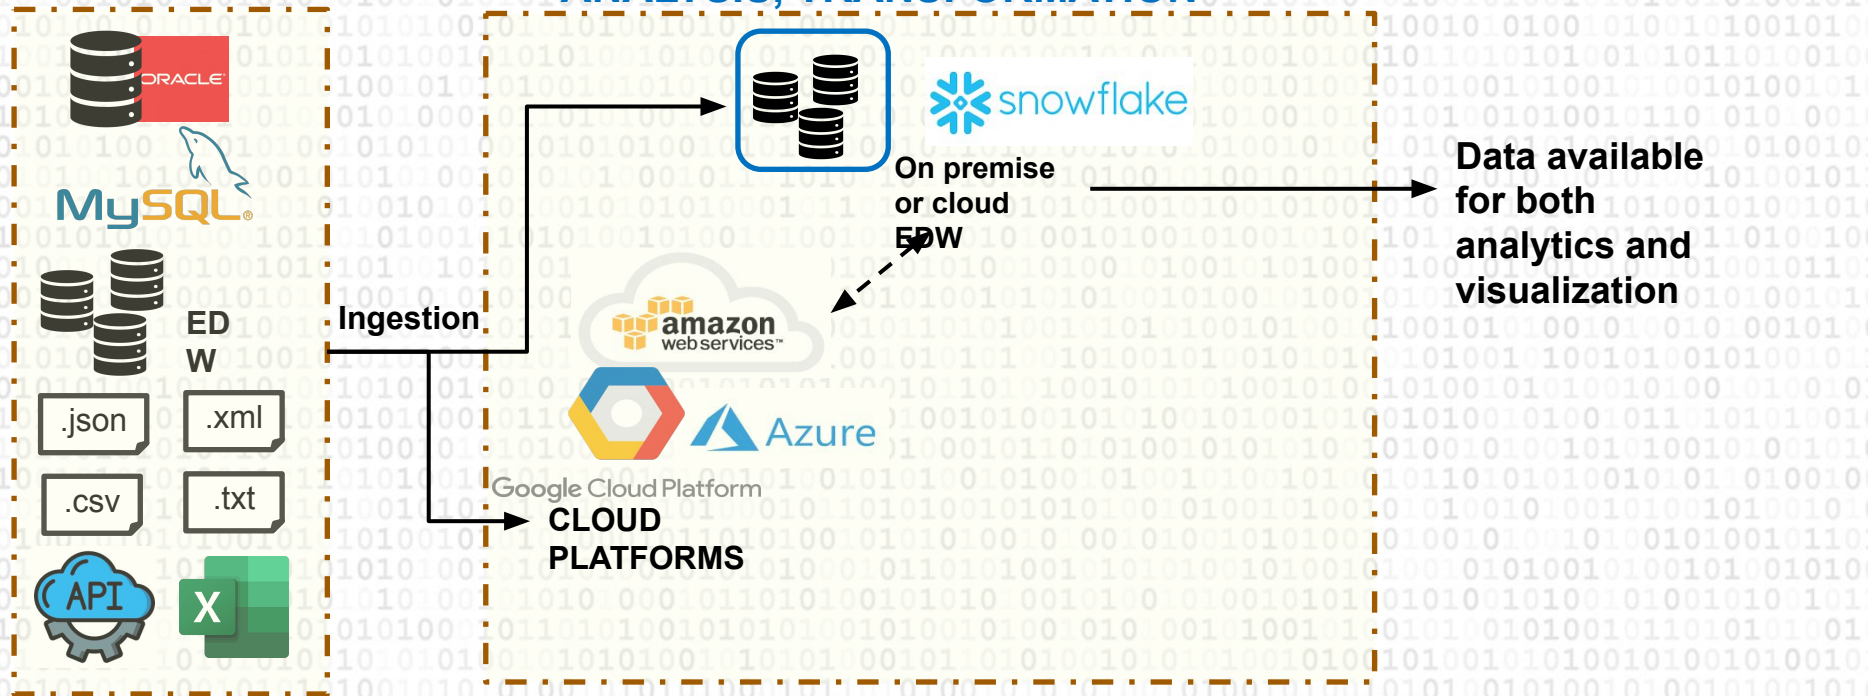


Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

# Data Architecture Diagram – Data Engineer

## DATA SOURCES

## DATA CLEANING, PROCESSING, ANALYSIS, TRANSFORMATION



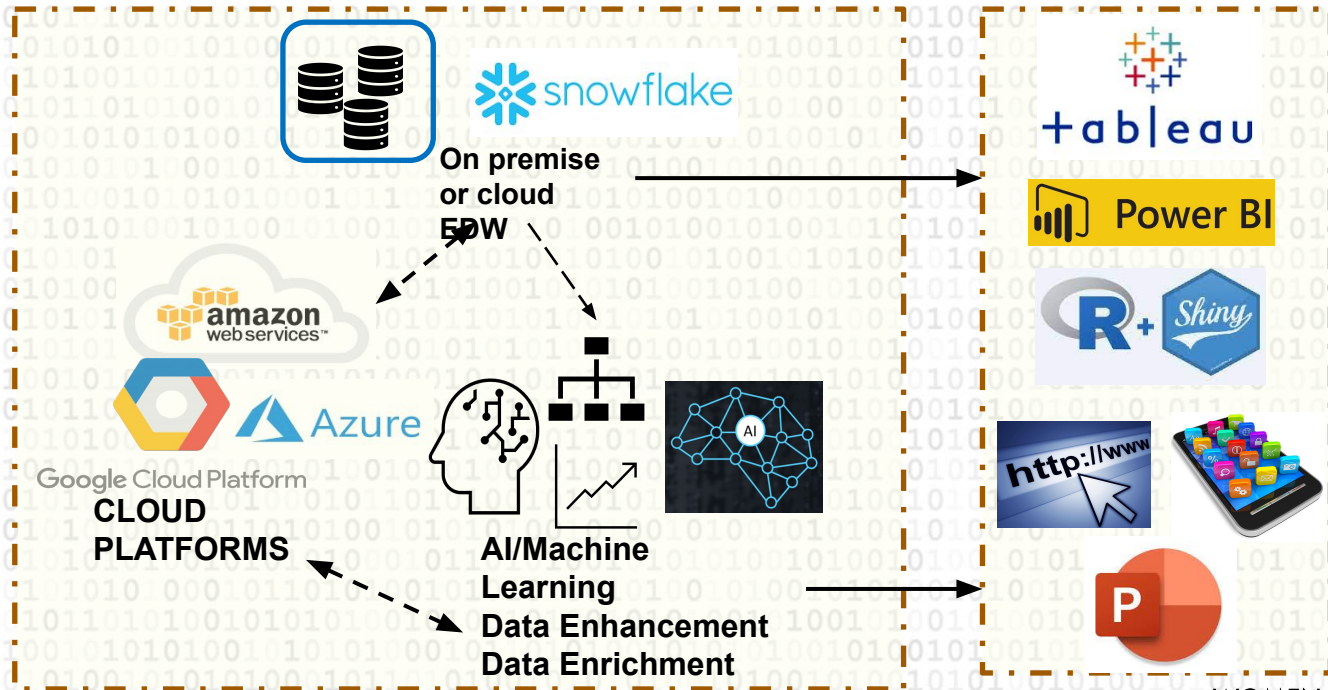
Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath



# Data Architecture Diagram – Data Scientist

**DATA CLEANING, PROCESSING,  
ANALYSIS, TRANSFORMATION**

**USER INTERFACE**



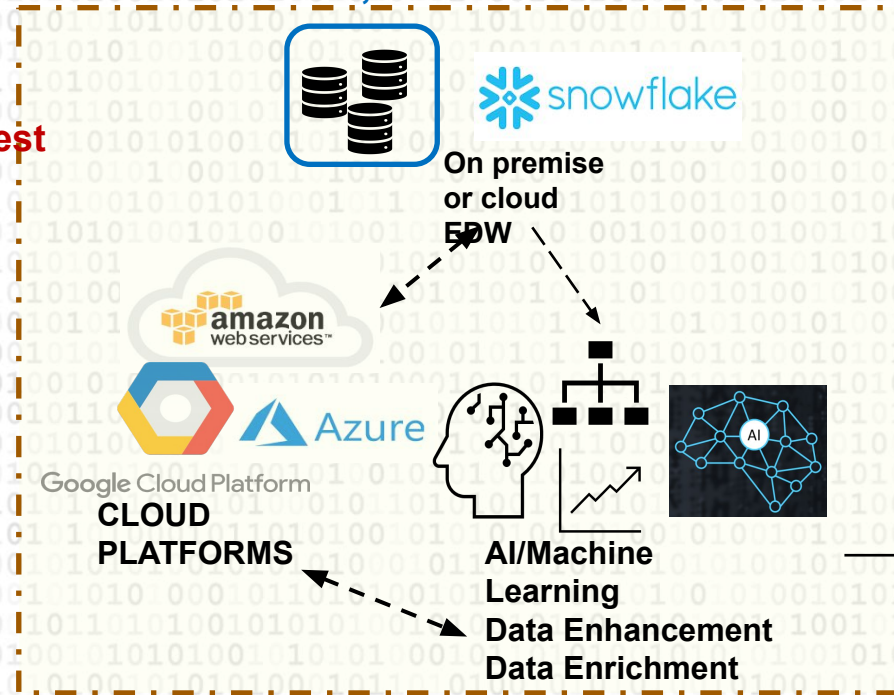
Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

WOMEN WHO  
**CODE**

# Data Architecture Diagram – ML Engineer

## DATA CLEANING, PROCESSING, ANALYSIS, TRANSFORMATION

**Automate and  
integrate with rest  
of the  
infrastructure**



Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

# Who are Data Engineers (DE), Data Scientists (DS) and Machine Learning Engineers (MLE)?

**You can unmute and talk or use the chat.**



# Data Engineer (DE) vs Data Scientist (DS) vs Machine Learning Engineer (MLE)

## **Data engineer:**

- Develops data ingestion pipelines
- Takes data from multiple sources and in multiple formats and combines to single source

## **Data scientist:**

- Finds patterns in the data to obtain insights
- Builds machine learning models to further enhance understanding of data

## **Machine learning engineer:**

- Develops scripts to integrate work of data scientists with the larger framework
- Automates the work of data scientist such that models can be triggered to run without a data scientist

# How do Data Engineers (DE), Data Scientists (DS) and Machine Learning Engineers (MLE) differ from each other?

**You can unmute and talk or use the chat.**

# How they are different/similar?

- Data quality analysis
- SQL queries

**DE**: Take data in a variety of formats, make it available in one queryable

- Integration of the pipelines with the overall architecture
- CI/CD

**DS**: Extensive data analysis, build ML models, perform statistical tests

**MLE**: Take ML models and deploy it using automated pipelines

Test different ML models

Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath



# NOTE

Data engineers, data scientists and machine learning engineers make use of existing/pre-built modules or libraries or functions.

They write their own functions or custom codes as well; however, not same as a software engineer or software developer.

Disclaimer: These can change based on companies/industry

Copyright © 2021 by Madhurima Nath

# A Day in a Life of Data Engineer

## Tasks:

- make available in one place datasets ingested from a variety of sources in a variety of formats
- build the data pipelines to ingest streaming and batch data from many sources
- pipelines perform extract, transform, and load (ETL) processes to make the data more usable
- convert from row-oriented formats to column-oriented formats to facilitate analytical queries, partition data, index it, tag it, govern it, etc.

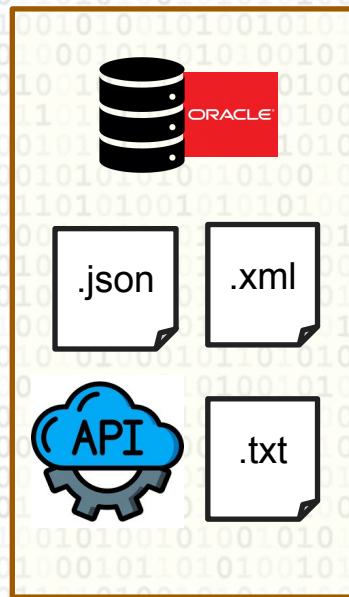
# A Day in a Life of Data Engineer

- make available in one place datasets ingested from a **variety of sources in a variety of formats**

## INTERNAL



## EXTERNAL



## INTERNAL



## manual file sources

\*API: Application Programming Interface: lets your product or service communicate with other products and services without having to know how they're implemented



# A Day in a Life of Data Engineer

- make available in one place datasets ingested from a **variety of sources in a variety of formats**



```
id name
age\n
01 ABCD
28\n
02 XYZ 39\n
...
...
```

```
id,name,age
01,ABCD,28
02,XYZ,39
...
...
```

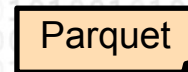
```
<log date=May 01, 2021>
  <message id=1>
    <person id=01></person>
    <text> hi, how are you
  </text>
    <time> May 01,
2021-16:09:22
    </time>
  </message>
...
...
```

```
{
  "text": "RT@abcland: my best day
...",
  "created_at": "Fri Apr 15 10:05:45
+0000 2020",
  "retweet_count": 0,
  "id_str": "561819028170009",
  "entities": {
    "user_mentions": [],
    "hashtags": []
  }, ...
}
```



# A Day in a Life of Data Engineer

- make available in one place datasets ingested from a **variety of sources in a variety of formats**



```
id name
age\n
01 ABCD
28\n
02 XYZ 39\n
...
```

```
id,name,age
01,ABCD,28
02,XYZ,39
...
```

```
<log date=May 01, 2021>
  <message id=1>
    <person id=01></person>
    <text> hi, how are you
  </text>
    <time> May 01,
2021-16:09:22
  </time>
</message>
```

```
{
  "text": "RT@abcland: my best day
  ..",
  "created_at": "Fri Apr 15 10:05:45
+0000 2020",
  "retweet_count": 0,
  "id_str": "561819028170009",
  "entities": {
    "user_mentions": [],
    "hashtags": []
  }, ...
}
```

## Can you identify which data formats are structured/unstructured/semi-structured?

You can unmute and talk or use the chat.

# A Day in a Life of Data Engineer

- make available in one place datasets ingested from a **variety of sources in a variety of formats**



```
id name
age\n
01 ABCD
28\n
02 XYZ 39\n
...
```

```
id,name,age
01,ABCD,28
02,XYZ,39
...
```

```
<log date=May 01, 2021>
  <message id=1>
    <person id=01></person>
    <text> hi, how are you
  </text>
    <time> May 01,
2021-16:09:22
    </time>
  </message>
...
```

```
{
  "text": "RT@abcland: my best day
  ",
  "created_at": "Fri Apr 15 10:05:45
+0000 2020",
  "retweet_count": 0,
  "id_str": "561819028170009",
  "entities": {
    "user_mentions": [],
    "hashtags": []
  }, ...
}
```



Parquet

UNSTRUCTURED

SEMI-STRUCTURED

STRUCTURED



# A Day in a Life of Data Engineer

- make available in one place datasets ingested from a **variety of sources in a variety of formats**



```
id name
age\n
01 ABCD
28\n
02 XYZ 39\n
...
```

```
id,name,age
01,ABCD,28
02,XYZ,39
...
```

```
<log date=May 01, 2021>
  <message id=1>
    <person id=01></person>
    <text> hi, how are you
  </text>
    <time> May 01,
2021-16:09:22
  </time>
</message>
```

```
{
  "text": "RT@abcland: my best day
  ..",
  "created_at": "Fri Apr 15 10:05:45
+0000 2020",
  "retweet_count": 0,
  "id_str": "561819028170009",
  "entities": {
    "user_mentions": [],
    "hashtags": []
  }, ...
}
```



## Which data formats are more user-friendly?

## Which data formats are more storage efficient?

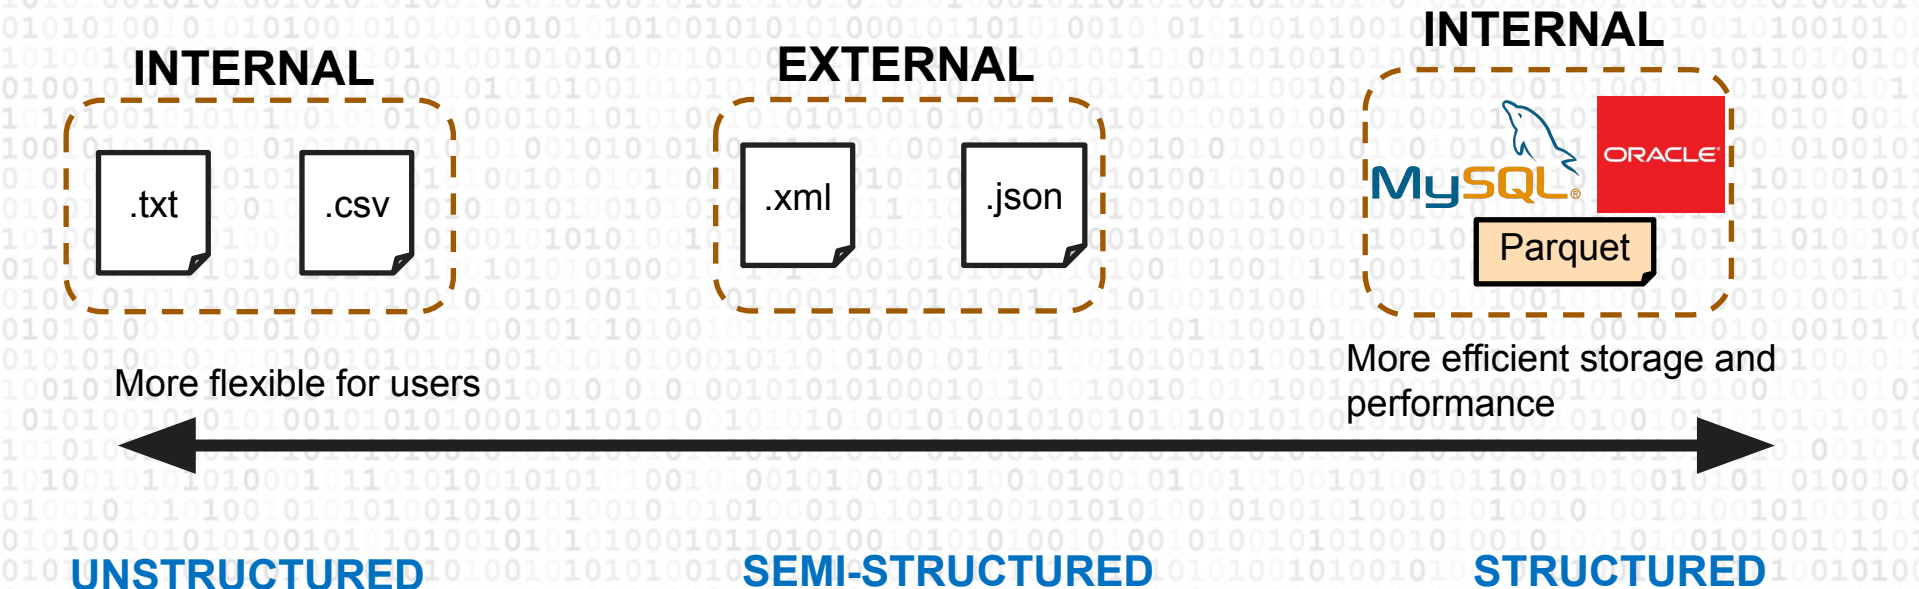
You can unmute and talk or use the chat.

Copyright © 2021 by Madhurima Nath

WOMEN WHO  
**CODE**

# A Day in a Life of Data Engineer

- make available in one place datasets ingested from a **variety of sources in a variety of formats**



Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

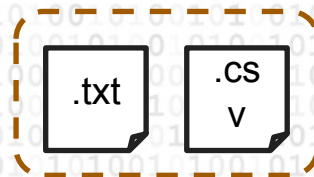
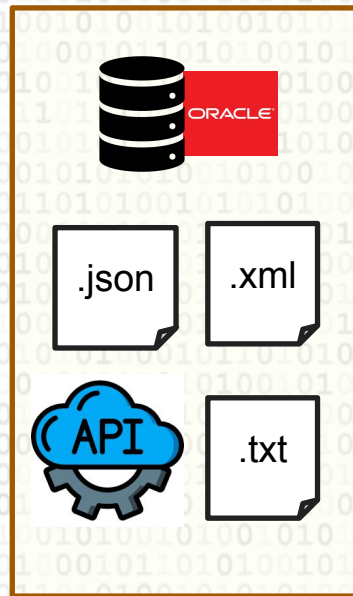
# A Day in a Life of Data Engineer

- Ingest data from internal and external sources – could be unstructured, semi-structured or structured and make it available at one place

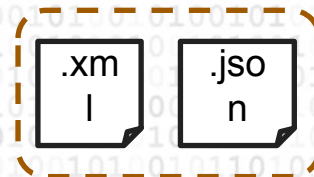
## INTERNAL



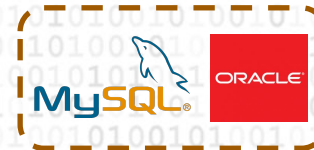
## EXTERNAL



UNSTRUCTURED



SEMI-STRUCTURED



STRUCTURED



# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming and batch data from many sources

## **Streaming data – real-time –**

e.g.: log files generated by customers using your mobile or web applications, ecommerce purchases, in-game player activity, information from social networks, financial trading floors, or geospatial services.

# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming and batch data from many sources

### **Streaming data – real-time –**

e.g.: log files generated by customers using your mobile or web applications, ecommerce purchases, in-game player activity, information from social networks, financial trading floors, or geospatial services.

```
from pyspark import SparkContext
from pyspark.streaming import StreamingContext
```

→ Calling pre-built  
spark functions

```
sc = SparkContext(master, appName)
ssc = StreamingContext(sc, 1)
```

“appName”: name for your application

“master”: Spark/local compute node

“1”: interval of 1 second



(<https://spark.apache.org/docs/latest/streaming-programming-guide.html>)

# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming and batch data from many sources

## **Batch data – not real-time –**

e.g.: very large amounts of data when data sources are legacy systems, moving data from databases/data warehouse to cloud.



# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming and batch data from many sources

### **Batch data – not real-time –**

e.g.: very large amounts of data when data sources are legacy systems, moving data from databases/data warehouse to cloud.

# read data in csv format

```
spark.read.format("csv")  
.option("mode", "FAILFAST")  
.option("inferSchema", "true")  
.option("path", "path/to/file(s)")  
.schema(someSchema)  
.load()
```

# write the data as parquet format

```
dataframe.write.format("parquet")  
.option("mode", "OVERWRITE")  
.option("dateFormat", "yyyy-MM-dd")  
.option("path", "path/to/file(s)")  
.save()
```

# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming (i.e., real-time) and batch data from many sources

## Writing files/data

```
dataframe.write.format("csv")  
.option("mode", "OVERWRITE")  
.option("dateFormat", "yyyy-MM-dd")  
.option("path", "path/to/file(s)")  
.save()
```

Save modes –

append: add output to already existing file at the same location

overwrite: completely overwrite over existing data/file

errorIfExists: throw an error and fail if data/file already exists at the same location

ignore: do nothing, if data/file already exists

# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming and batch data from many sources

## Streaming data – real-time –

e.g.: log files generated by customers using your mobile or web applications, ecommerce purchases, in-game player activity, information from social networks, financial trading floors, or geospatial services.



databricks

# read data continuously from AWS S3 location

```
val inputDF = spark.readStream.json("s3://logs")
```

([https://databricks.com/blog/2016/07/28/structured-streaming-in-  
apache-spark.html](https://databricks.com/blog/2016/07/28/structured-streaming-in-apache-spark.html))

# write the data into MySQL database

```
inputDF.groupBy($"action", window($"time", "1 hour")).count()  
  .writeStream.format("jdbc")  
  .start("jdbc:mysql://...")
```

Disclaimer: These can change based on companies/industry

Copyright © 2021 by Madhurima Nath

WOMEN WHO  
**CODE**



# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming and batch data from many sources

## Batch data – not real-time –

e.g.: very large amounts of data when data sources are legacy systems, moving data from databases/data warehouse to cloud.



databricks

(<https://databricks.com/blog/2016/07/28/structured-streaming-in-apache-spark.html>)

# read data in JSON format from AWS S3 location

```
val inputDF = spark.read.json("s3://logs")
```



# write the data into MySQL database

```
inputDF.groupBy($"action", window($"time", "1 hour")).count()  
  .write.format("jdbc")  
  .start("jdbc:mysql://...")
```

Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

WOMEN WHO  
**CODE**

# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming and batch data from many sources

**Can you give examples of streaming and batch data?**

**You can unmute and talk or use the chat.**

# A Day in a Life of Data Engineer

## Tasks:

- build the data pipelines to ingest streaming and batch data from many sources

## Examples of streaming data

- e-commerce data
- social media data
- financial trading data

## Examples of streaming data

- data from Oracle/Postgres databases
- data from enterprise data warehouses (EDW)
- data from enterprise resource planning (ERP)

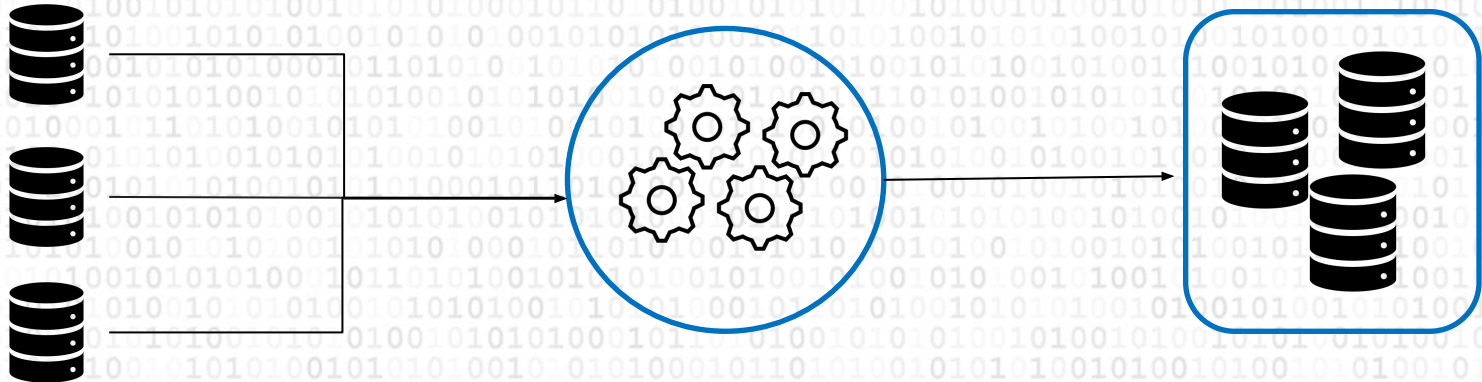


# A Day in a Life of Data Engineer

## Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes to make the data more usable

- Extract (E) – data from sources (databases)
- Transform (T) – data to match certain common defined format (for business purposes)
- Load (L) – load re-formatted data to data warehouse



# A Day in a Life of Data Engineer

Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



Example:

Suppose you want to extract data from multiple databases and put it in cloud data storage.

## How would you extract the data to load into cloud data storage?

**You can unmute and talk or use the chat.**

Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

# A Day in a Life of Data Engineer

Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



Example:

Suppose you want to extract data from multiple databases and put it in cloud data storage.

Step 1: How to handle day 0 load (everything that currently exist)?

Step 2: How to handle new incremental data?



# A Day in a Life of Data Engineer

Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



Example:

Suppose you want to extract data from multiple databases and put it in cloud data storage.

Step 1: Day 0 (already existing data) load

- Partition the data
- Keep track of the partitions
- Load the partitioned data

# A Day in a Life of Data Engineer

Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



Example:

Suppose you want to extract data from multiple databases and put it in cloud data storage.

Step 1: Day 0 (already existing data) load

Step 2: New incremental data load

- decide on the frequency (daily/weekly/monthly)
- refresh the tables with new data

Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

# A Day in a Life of Data Engineer

Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes

Example:

Suppose you want to extract data from multiple databases and put it in cloud data storage. (Code for Azure data storage on Databricks)

```
table_id = dbutils.widgets.get("table_id")
job_id = dbutils.widgets.get("job_id")
```

```
storageCfg = DatalakeStorageConfig()
storageCfgPath = ConfigLoader(DatalakeStorageConfig().cfg_path)
```

- ☐ get source table id from source
- ☐ get job id
- ☐ load configuration of cloud storage
- ☐ load configuration path of cloud storage

Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath





# A Day in a Life of Data Engineer

Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes

Example:

Suppose you want to extract data from multiple databases and put it in cloud data storage. (Code for Azure data storage on Databricks)



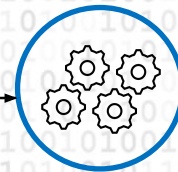
```
table_name = storageCfgPath.get_dataset_by_id(table_id).table_name
table_folder = storageCfgPath.get_dataset_by_id(table_id).folder_name
source_path = f'../{storageCfg.root_dir}/../{table_folder}'
dest_path = f'../{storageCfg.root_dir}/../{table_folder}'
```

- ☐ set destination table name
- ☐ set destination table folder
- ☐ set source table path as parametrized path
- ☐ set target table path as parametrized path

# A Day in a Life of Data Engineer

## Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



```
import json
from pyspark.sql.types import *
from pyspark.sql import functions as F
```

Call pre-built python  
and spark functions

```
source_df = spark.read.format('parquet')
                .load(source_path).filter(F.col('job_id') == job_id)

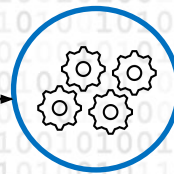
if source_df.count() > 0:
    ( source_df.write.format('format').mode('merge').partitionBy('job_id').save(source_path) )
try:
    sqlCmd = 'CREATE TABLE IF NOT EXISTS table' + table_name + " USING LOCATION '{ }'"
                .format(dest__path)
    spark.sql(sqlCmd)
except:
    OSError as e
```

Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

# A Day in a Life of Data Engineer

## Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



import json

from pyspark.sql.types import \*

from pyspark.sql import functions as F

→ Call pre-built python and spark functions

```
source_df = spark.read.format('parquet')  
.load(source_path).filter(F.col('job_id') == job_id)
```

→ Read batch files (format – parquet)  
Filter these files with job id

```
if source_df.count() > 0:
```

```
    (source_df.write.format('format').mode('merge').partitionBy('job_id').save(source_path))
```

```
    try:
```

```
        sqlCmd = 'CREATE TABLE IF NOT EXISTS table' + table_name + " USING LOCATION '{ }'"  
        .format(dest_path)
```

```
        spark.sql(sqlCmd)
```

```
    except:
```

```
        OSError as e
```

Disclaimer: These can change based on companies/industry

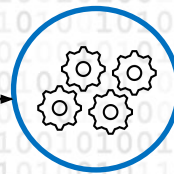
Copyright © 2021 by Madhurima Nath



# A Day in a Life of Data Engineer

## Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



import json

from pyspark.sql.types import \*

from pyspark.sql import functions as F

Call pre-built python  
and spark functions

source\_df = spark.read.format('parquet')

.load(source\_path).filter(F.col('job\_id') == job\_id)

Read batch files (format –  
parquet)  
Filter these files with job id

```
if source_df.count() > 0:
```

```
    ( source_df.write.format('format').mode('merge').partitionBy('job_id').save(source_path) )
```

```
    try:
```

```
        sqlCmd = 'CREATE TABLE IF NOT EXISTS table' + table_name + " USING LOCATION '{ }'"
```

```
        .format(dest__path)
```

```
        spark.sql(sqlCmd)
```

```
    except:
```

```
        OSError as e
```

Save data to  
defined path  
Create table if  
table doesn't  
exist at defined  
location

Disclaimer: These can change based on companies/industry

Copyright © 2021 by Madhurima Nath

WOMEN WHO  
**CODE**

# A Day in a Life of Data Engineer

Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



## How would you handle streaming data? What would be the steps?

**You can unmute and talk or use the chat.**

Disclaimer: These can change based on companies/industry  
Copyright © 2021 by Madhurima Nath

# A Day in a Life of Data Engineer

Tasks:

- pipelines perform **Extract, Transform, and Load (ETL)** processes



Streaming data

Step 1: Day 0 load – historical data load

- partition the data
- load the partitioned data

Step 2: Real-time incremental data

- decide on the frequency of refresh – seconds/minutes/hours
- refresh data in the tables



# Tech stack for data engineers

Languages: **SQL**, **NoSQL**, **Python**, **Java**, **Scala**  
**Spark**, **Databricks**, **Kafka**, **Hadoop**, **CI/CD\***  
framework



**nosql**



**databricks**



\*CI/CD: Continuous Integration Continuous Deployment

Disclaimer: These can change based on companies/industry

# Backend Study Group

- WWCode [Presentation](#) and [Demo](#)
- [WWCode YouTube channel](#) for session recordings.

## Resources:

[spark documentation](#)

[databricks documentation](#)

[snowflake documentation](#)

[azure data engineer certification course](#)

[google cloud data engineer certification course](#)

[aws big data certification](#)

[Udacity data engineering interview questions](#)



# Q+A

WOMEN WHO  
**CODE**